

Class 13: RNA seq analysis with DESeq2

Hsiang-Ying Lu (PID: A15608316)

The data for this hands-on session comes from a published RNA-seq experiment where airway smooth muscle cells were treated with dexamethasone, a synthetic glucocorticoid steroid with anti-inflammatory effects (Himes et al. 2014).

```
library(DESeq2)
```

Data import

```
# Complete the missing code
counts <- read.csv("airway_scaledcounts.csv", row.names=1)
metadata <- read.csv("airway_metadata.csv")
```

```
head(counts)
```

| | SRR1039508 | SRR1039509 | SRR1039512 | SRR1039513 | SRR1039516 |
|------------------|------------|------------|------------|------------|------------|
| ENSG00000000003 | 723 | 486 | 904 | 445 | 1170 |
| ENSG00000000005 | 0 | 0 | 0 | 0 | 0 |
| ENSG000000000419 | 467 | 523 | 616 | 371 | 582 |
| ENSG000000000457 | 347 | 258 | 364 | 237 | 318 |
| ENSG000000000460 | 96 | 81 | 73 | 66 | 118 |
| ENSG000000000938 | 0 | 0 | 1 | 0 | 2 |

| | SRR1039517 | SRR1039520 | SRR1039521 |
|------------------|------------|------------|------------|
| ENSG00000000003 | 1097 | 806 | 604 |
| ENSG00000000005 | 0 | 0 | 0 |
| ENSG000000000419 | 781 | 417 | 509 |
| ENSG000000000457 | 447 | 330 | 324 |
| ENSG000000000460 | 94 | 102 | 74 |
| ENSG000000000938 | 0 | 0 | 0 |

```
dim(counts)
```

```
[1] 38694      8
```

```
head(metadata)
```

| | id | dex | celltype | geo_id |
|---|------------|---------|----------|------------|
| 1 | SRR1039508 | control | N61311 | GSM1275862 |
| 2 | SRR1039509 | treated | N61311 | GSM1275863 |
| 3 | SRR1039512 | control | N052611 | GSM1275866 |
| 4 | SRR1039513 | treated | N052611 | GSM1275867 |
| 5 | SRR1039516 | control | N080611 | GSM1275870 |
| 6 | SRR1039517 | treated | N080611 | GSM1275871 |

```
sum(metadata$dex == "control")
```

```
[1] 4
```

```
table(metadata$dex)
```

| control | treated |
|---------|---------|
| 4 | 4 |

Q1. How many genes are in this dataset? 38694

Q2. How many 'control' cell lines do we have? 4

I want to compare the control to the treated columns. To do this I will

- Step 1. Identify and extract the "control" columns.
- Step 2. Calculate the mean value per gene for all these "control" columns and save as `control.mean`.
- Step 3. Do the same for treated
- Step 4. Compare the `control.mean` and `treated.mean` values.

Step 1:

```
control.inds <- metadata$dex=="control"
```

```
head(counts[,control.inds])
```

| | SRR1039508 | SRR1039512 | SRR1039516 | SRR1039520 |
|------------------|------------|------------|------------|------------|
| ENSG000000000003 | 723 | 904 | 1170 | 806 |
| ENSG000000000005 | 0 | 0 | 0 | 0 |
| ENSG000000000419 | 467 | 616 | 582 | 417 |
| ENSG000000000457 | 347 | 364 | 318 | 330 |
| ENSG000000000460 | 96 | 73 | 118 | 102 |
| ENSG000000000938 | 0 | 1 | 2 | 0 |

```
control.means <- rowMeans(counts[,control.inds])
head(control.means)
```

| ENSG000000000003 | ENSG000000000005 | ENSG000000000419 | ENSG000000000457 | ENSG000000000460 |
|------------------|------------------|------------------|------------------|------------------|
| 900.75 | 0.00 | 520.50 | 339.75 | 97.25 |
| ENSG000000000938 | | | | |
| 0.75 | | | | |

```
treated.inds <- metadata$dex=="treated"
```

```
head(counts[,treated.inds])
```

| | SRR1039509 | SRR1039513 | SRR1039517 | SRR1039521 |
|------------------|------------|------------|------------|------------|
| ENSG000000000003 | 486 | 445 | 1097 | 604 |
| ENSG000000000005 | 0 | 0 | 0 | 0 |
| ENSG000000000419 | 523 | 371 | 781 | 509 |
| ENSG000000000457 | 258 | 237 | 447 | 324 |
| ENSG000000000460 | 81 | 66 | 94 | 74 |
| ENSG000000000938 | 0 | 0 | 0 | 0 |

```
treated.means <- rowMeans(counts[,treated.inds])
head(treated.means)
```

| ENSG000000000003 | ENSG000000000005 | ENSG000000000419 | ENSG000000000457 | ENSG000000000460 |
|------------------|------------------|------------------|------------------|------------------|
| 658.00 | 0.00 | 546.00 | 316.50 | 78.75 |
| ENSG000000000938 | | | | |
| 0.00 | | | | |


We will combine our meancount data for bookkeeping purpose

```
meancounts <- data.frame(control.means, treated.means)
colSums(meancounts)
```

```
control.means treated.means
23005324      22196524
```


Let's see what these count values look like...

```
plot(meancounts)
```



```
library(ggplot2)


ggplot(meancounts)+
  aes(control.means, treated.means)+
  geom_point(alpha=0.2)
```



```
plot(meancounts, log="xy")
```

Warning in xy.coords(x, y, xlabel, ylabel, log): 15032 x values ≤ 0 omitted from logarithmic plot

Warning in xy.coords(x, y, xlabel, ylabel, log): 15281 y values ≤ 0 omitted from logarithmic plot



Logs are super useful when we have such skewed data

```
#Treated/control
log2(20/10)
```

```
[1] 1
```

Add log2(Fold-change) values to our wee results table.

```
meancounts$log2fc <- log2(meancounts$treated.means/meancounts$control.means)
head(meancounts)
```

| | control.means | treated.means | log2fc |
|------------------|---------------|---------------|-------------|
| ENSG000000000003 | 900.75 | 658.00 | -0.45303916 |
| ENSG000000000005 | 0.00 | 0.00 | NaN |
| ENSG000000000419 | 520.50 | 546.00 | 0.06900279 |
| ENSG000000000457 | 339.75 | 316.50 | -0.10226805 |
| ENSG000000000460 | 97.25 | 78.75 | -0.30441833 |
| ENSG000000000938 | 0.75 | 0.00 | -Inf |

I need to exclude any genes with zero counts as we can't say anything about them anyway from this experiment and it causes me math pain.

```
# What values in the first two cols are zero

to.rm.inds <- rowSums(meancounts[,1:2] == 0)>0
mycounts <- meancounts[!to.rm.inds, ]
```

Q. How many genes do I have left?

```
nrow(mycounts)
```

```
[1] 21817
```

Q. How many genes are “up-degulated” i.e. have a log2(fold-change) greater than +2?

```
sum(mycounts$log2fc > +2)
```

```
[1] 250
```

Q. How many are “down-regulated” with a log2(fold-change) less than -2?

```
sum(mycounts$log2fc < -2)
```

```
[1] 367
```

Q10. Do you trust these results? Why or why not? No, because there's no information on statistical significance.

Running DESeq

Like many bioconductor analysis packages DESeq wants its input in a very particular way.

```
dds <- DESeqDataSetFromMatrix(countData=counts,
                              colData = metadata,
                              design =~dex)
```

converting counts to integer mode

Warning in DESeqDataSet(se, design = design, ignoreRank): some variables in design formula are characters, converting to factors

To run DESeq analysis we call the main function from the package called DESeq(dds)

```
dds <- DESeq(dds)
```

estimating size factors

estimating dispersions

gene-wise dispersion estimates

mean-dispersion relationship

final dispersion estimates

fitting model and testing

To get the results out of this dds object we can use the DESeq results() function.

```
res <- results(dds)
head(res)
```

log2 fold change (MLE): dex treated vs control

Wald test p-value: dex treated vs control

DataFrame with 6 rows and 6 columns

| | baseMean | log2FoldChange | lfcSE | stat | pvalue |
|------------------|------------|----------------|-----------|-----------|-----------|
| | <numeric> | <numeric> | <numeric> | <numeric> | <numeric> |
| ENSG000000000003 | 747.194195 | -0.3507030 | 0.168246 | -2.084470 | 0.0371175 |
| ENSG000000000005 | 0.000000 | NA | NA | NA | NA |
| ENSG000000000419 | 520.134160 | 0.2061078 | 0.101059 | 2.039475 | 0.0414026 |
| ENSG000000000457 | 322.664844 | 0.0245269 | 0.145145 | 0.168982 | 0.8658106 |
| ENSG000000000460 | 87.682625 | -0.1471420 | 0.257007 | -0.572521 | 0.5669691 |
| ENSG000000000938 | 0.319167 | -1.7322890 | 3.493601 | -0.495846 | 0.6200029 |
| | padj | | | | |
| | <numeric> | | | | |
| ENSG000000000003 | 0.163035 | | | | |
| ENSG000000000005 | NA | | | | |


```

ENSG000000000419 0.176032
ENSG000000000457 0.961694
ENSG000000000460 0.815849
ENSG000000000938 NA


```

A common summary visualization is called a Volcano plot.

```

plot(res$log2FoldChange, -log(res$padj),
     xlab="Log2 Fold-Change",
     ylab="-log p-value")
abline(v=c(-2, 2), col="red")
abline(h=-log(0.05), col="blue")

```




```

mycols <- rep("gray", nrow(res))
mycols[res$log2FoldChange > 2] <- "black"
mycols[res$log2FoldChange < -2] <- "black"
mycols[res$padj > 0.05] <- "gray"

plot(res$log2FoldChange, -log(res$padj), col=mycols,
     xlab="Log2 Fold-Change",
     ylab="-log p-value")

```

```
abline(v=c(-2, 2), col="red")
abline(h=-log(0.05), col="blue")
```



Save our results to date

```
write.csv(res, file="myresults.csv")
```

adding annotation data

We need to translate or “map” our ensemble IDs into more understandable gene names

```
library("AnnotationDbi")
library("org.Hs.eg.db")
```

```
columns(org.Hs.eg.db)
```

```
[1] "ACCNUM"      "ALIAS"      "ENSEMBL"    "ENSEMBLPROT" "ENSEMBLTRANS"
[6] "ENTREZID"    "ENZYME"     "EVIDENCE"   "EVIDENCEALL" "GENENAME"
[11] "GENETYPE"    "GO"         "GOALL"      "IPI"         "MAP"
[16] "OMIM"        "ONTOLOGY"   "ONTOLOGYALL" "PATH"        "PFAM"
[21] "PMID"        "PROSITE"    "REFSEQ"     "SYMBOL"      "UCSCKG"
[26] "UNIPROT"
```

```
res$symbol <- mapIds(org.Hs.eg.db,
  keys=row.names(res), # Our genenames
  keytype="ENSEMBL",   # The format of our genenames
  column="SYMBOL",     # The new format we want to add
  multiVals="first")
```

'select()' returned 1:many mapping between keys and columns

```
head(res)
```

log2 fold change (MLE): dex treated vs control

Wald test p-value: dex treated vs control

DataFrame with 6 rows and 7 columns

| | baseMean | log2FoldChange | lfcSE | stat | pvalue |
|-------------------|------------|----------------|-----------|-----------|-----------|
| | <numeric> | <numeric> | <numeric> | <numeric> | <numeric> |
| ENSG000000000003 | 747.194195 | -0.3507030 | 0.168246 | -2.084470 | 0.0371175 |
| ENSG000000000005 | 0.000000 | NA | NA | NA | NA |
| ENSG0000000000419 | 520.134160 | 0.2061078 | 0.101059 | 2.039475 | 0.0414026 |
| ENSG0000000000457 | 322.664844 | 0.0245269 | 0.145145 | 0.168982 | 0.8658106 |
| ENSG0000000000460 | 87.682625 | -0.1471420 | 0.257007 | -0.572521 | 0.5669691 |
| ENSG0000000000938 | 0.319167 | -1.7322890 | 3.493601 | -0.495846 | 0.6200029 |

| | padj | symbol |
|-------------------|-----------|-------------|
| | <numeric> | <character> |
| ENSG0000000000003 | 0.163035 | TSPAN6 |
| ENSG0000000000005 | NA | TNMD |
| ENSG0000000000419 | 0.176032 | DPM1 |
| ENSG0000000000457 | 0.961694 | SCYL3 |
| ENSG0000000000460 | 0.815849 | FIRRM |
| ENSG0000000000938 | NA | FGR |

```
res$entrez <- mapIds(org.Hs.eg.db,
  keys=row.names(res), # Our genenames
```

```
keytype="ENSEMBL", # The format of our genenames
column="ENTREZID", # The new format we want to add
multiVals="first")
```

'select()' returned 1:many mapping between keys and columns

```
head(res)
```

log2 fold change (MLE): dex treated vs control

Wald test p-value: dex treated vs control

DataFrame with 6 rows and 8 columns

| | baseMean | log2FoldChange | lfcSE | stat | pvalue |
|------------------|------------|----------------|-------------|-----------|-----------|
| | <numeric> | <numeric> | <numeric> | <numeric> | <numeric> |
| ENSG000000000003 | 747.194195 | -0.3507030 | 0.168246 | -2.084470 | 0.0371175 |
| ENSG000000000005 | 0.000000 | NA | NA | NA | NA |
| ENSG000000000419 | 520.134160 | 0.2061078 | 0.101059 | 2.039475 | 0.0414026 |
| ENSG000000000457 | 322.664844 | 0.0245269 | 0.145145 | 0.168982 | 0.8658106 |
| ENSG000000000460 | 87.682625 | -0.1471420 | 0.257007 | -0.572521 | 0.5669691 |
| ENSG000000000938 | 0.319167 | -1.7322890 | 3.493601 | -0.495846 | 0.6200029 |
| | padj | symbol | entrez | | |
| | <numeric> | <character> | <character> | | |
| ENSG000000000003 | 0.163035 | TSPAN6 | 7105 | | |
| ENSG000000000005 | NA | TNMD | 64102 | | |
| ENSG000000000419 | 0.176032 | DPM1 | 8813 | | |
| ENSG000000000457 | 0.961694 | SCYL3 | 57147 | | |
| ENSG000000000460 | 0.815849 | FIRRM | 55732 | | |
| ENSG000000000938 | NA | FGR | 2268 | | |

```
res$uniprot <- mapIds(org.Hs.eg.db,
                      keys=row.names(res), # Our genenames
                      keytype="ENSEMBL", # The format of our genenames
                      column="UNIPROT", # The new format we want to add
                      multiVals="first")
```

'select()' returned 1:many mapping between keys and columns

```
head(res)
```

log2 fold change (MLE): dex treated vs control

Wald test p-value: dex treated vs control

DataFrame with 6 rows and 9 columns

| | baseMean | log2FoldChange | lfcSE | stat | pvalue |
|------------------|------------|----------------|-------------|-------------|-----------|
| | <numeric> | <numeric> | <numeric> | <numeric> | <numeric> |
| ENSG000000000003 | 747.194195 | -0.3507030 | 0.168246 | -2.084470 | 0.0371175 |
| ENSG000000000005 | 0.000000 | NA | NA | NA | NA |
| ENSG000000000419 | 520.134160 | 0.2061078 | 0.101059 | 2.039475 | 0.0414026 |
| ENSG000000000457 | 322.664844 | 0.0245269 | 0.145145 | 0.168982 | 0.8658106 |
| ENSG000000000460 | 87.682625 | -0.1471420 | 0.257007 | -0.572521 | 0.5669691 |
| ENSG000000000938 | 0.319167 | -1.7322890 | 3.493601 | -0.495846 | 0.6200029 |
| | padj | symbol | entrez | uniprot | |
| | <numeric> | <character> | <character> | <character> | |
| ENSG000000000003 | 0.163035 | TSPAN6 | 7105 | AOA024RCIO | |
| ENSG000000000005 | NA | TNMD | 64102 | Q9H2S6 | |
| ENSG000000000419 | 0.176032 | DPM1 | 8813 | O60762 | |
| ENSG000000000457 | 0.961694 | SCYL3 | 57147 | Q8IZE3 | |
| ENSG000000000460 | 0.815849 | FIRRM | 55732 | AOA024R922 | |
| ENSG000000000938 | NA | FGR | 2268 | P09769 | |

```
res$genenames <- mapIds(org.Hs.eg.db,  
  keys=row.names(res), # Our genenames  
  keytype="ENSEMBL",   # The format of our genenames  
  column="GENENAME",   # The new format we want to add  
  multiVals="first")
```

'select()' returned 1:many mapping between keys and columns

```
head(res)
```

log2 fold change (MLE): dex treated vs control

Wald test p-value: dex treated vs control

DataFrame with 6 rows and 10 columns

| | baseMean | log2FoldChange | lfcSE | stat | pvalue |
|------------------|------------|----------------|-----------|-----------|-----------|
| | <numeric> | <numeric> | <numeric> | <numeric> | <numeric> |
| ENSG000000000003 | 747.194195 | -0.3507030 | 0.168246 | -2.084470 | 0.0371175 |
| ENSG000000000005 | 0.000000 | NA | NA | NA | NA |
| ENSG000000000419 | 520.134160 | 0.2061078 | 0.101059 | 2.039475 | 0.0414026 |
| ENSG000000000457 | 322.664844 | 0.0245269 | 0.145145 | 0.168982 | 0.8658106 |
| ENSG000000000460 | 87.682625 | -0.1471420 | 0.257007 | -0.572521 | 0.5669691 |

| | padj | symbol | entrez | uniprot |
|------------------|-----------|-------------|-------------|-------------|
| | <numeric> | <character> | <character> | <character> |
| ENSG000000000938 | 0.319167 | -1.7322890 | 3.493601 | -0.495846 |
| | | | | 0.6200029 |
| ENSG000000000003 | 0.163035 | TSPAN6 | 7105 | AOA024RCIO |
| ENSG000000000005 | NA | TNMD | 64102 | Q9H2S6 |
| ENSG000000000419 | 0.176032 | DPM1 | 8813 | 060762 |
| ENSG000000000457 | 0.961694 | SCYL3 | 57147 | Q8IZE3 |
| ENSG000000000460 | 0.815849 | FIRRM | 55732 | AOA024R922 |
| ENSG000000000938 | NA | FGR | 2268 | P09769 |

| | genenames |
|------------------|------------------------|
| | <character> |
| ENSG000000000003 | tetraspanin 6 |
| ENSG000000000005 | tenomodulin |
| ENSG000000000419 | dolichyl-phosphate m.. |
| ENSG000000000457 | SCY1 like pseudokina.. |
| ENSG000000000460 | FIGNL1 interacting r.. |
| ENSG000000000938 | FGR proto-oncogene, .. |

Pathway analysis

```
library(pathview)
```

```
#####
Pathview is an open source software package distributed under GNU General
Public License version 3 (GPLv3). Details of GPLv3 is available at
http://www.gnu.org/licenses/gpl-3.0.html. Particullary, users are required to
formally cite the original Pathview paper (not just mention it) in publications
or products. For details, do citation("pathview") within R.
```

```
The pathview downloads and uses KEGG data. Non-academic uses may require a KEGG
license agreement (details at http://www.kegg.jp/kegg/legal.html).
```

```
#####
```

```
library(gage)
```

```
library(gageData)

data(kegg.sets.hs)

# Examine the first 2 pathways in this kegg set for humans
head(kegg.sets.hs, 2)
```

```
$`hsa00232 Caffeine metabolism`
```

```
[1] "10" "1544" "1548" "1549" "1553" "7498" "9"
```

```
$`hsa00983 Drug metabolism - other enzymes`
```

```
[1] "10" "1066" "10720" "10941" "151531" "1548" "1549" "1551"
[9] "1553" "1576" "1577" "1806" "1807" "1890" "221223" "2990"
[17] "3251" "3614" "3615" "3704" "51733" "54490" "54575" "54576"
[25] "54577" "54578" "54579" "54600" "54657" "54658" "54659" "54963"
[33] "574537" "64816" "7083" "7084" "7172" "7363" "7364" "7365"
[41] "7366" "7367" "7371" "7372" "7378" "7498" "79799" "83549"
[49] "8824" "8833" "9" "978"
```

```
foldchanges = res$log2FoldChange
names(foldchanges) = res$entrez
head(foldchanges)
```

```
7105 64102 8813 57147 55732 2268
-0.35070302 NA 0.20610777 0.02452695 -0.14714205 -1.73228897
```

Run gage:

```
# Get the results
keggres = gage(foldchanges, gsets=kegg.sets.hs)
```

```
attributes(keggres)
```

```
$names
```

```
[1] "greater" "less" "stats"
```

```
# Look at the first three down (less) pathways
head(keggres$less, 3)
```

| | | p.geomean | stat.mean | p.val |
|----------|---------------------------|--------------|-----------|--------------|
| hsa05332 | Graft-versus-host disease | 0.0004250461 | -3.473346 | 0.0004250461 |
| hsa04940 | Type I diabetes mellitus | 0.0017820293 | -3.002352 | 0.0017820293 |
| hsa05310 | Asthma | 0.0020045888 | -3.009050 | 0.0020045888 |

| | | q.val | set.size | exp1 |
|----------|---------------------------|------------|----------|--------------|
| hsa05332 | Graft-versus-host disease | 0.09053483 | 40 | 0.0004250461 |
| hsa04940 | Type I diabetes mellitus | 0.14232581 | 42 | 0.0017820293 |
| hsa05310 | Asthma | 0.14232581 | 29 | 0.0020045888 |

Let's have a look at one of these pathways

```
pathview(gene.data=foldchanges, pathway.id="hsa05310")
```

'select()' returned 1:1 mapping between keys and columns

Info: Working in directory /Users/sophialu1999/Desktop/UCSD Biological Sciences Ph.D./BGGN21

Info: Writing image file hsa05310.pathview.png

